Technical Analysis and Strategic Implementation Framework for Advanced Image Tampering Detection and Localization

The proliferation of sophisticated digital image editing tools and the advent of generative artificial intelligence have rendered the task of distinguishing between authentic and manipulated visual media increasingly complex. For digital forensics, the challenge is no longer merely identifying visual anomalies but detecting subtle statistical inconsistencies that occur at the sub-pixel level. This report provides an exhaustive technical analysis of the state-of-the-art in image manipulation detection and localization (IMDL), focusing on the integration of artifact extraction, robust mathematical optimization, and performance-oriented deployment within resource-constrained environments such as Google Colab.

Theoretical Foundations of Digital Image Forensics

Image forensics is broadly categorized into active and passive methodologies. Active methods, such as digital watermarking and signatures, rely on the presence of embedded information to verify authenticity.[1] However, these are often impractical for "in-the-wild" scenarios where the original image is unavailable. Passive methods, therefore, represent the primary research frontier, focusing on the intrinsic properties of the digital image itself.[2] Passive forgery detection aims to identify specific classes of manipulation: splicing, where a region from one image is pasted into another; copy-move, where a region is duplicated within the same image; and removal or inpainting, where an object is deleted and the area is filled using generative techniques.[4]

The core assumption in digital forensics is that every stage of the image acquisition pipeline—from the light hitting the sensor to the final JPEG compression—leaves a specific digital footprint. Manipulations inevitably disrupt these footprints, creating inconsistencies in noise patterns, compression artifacts, and pixel correlations.[7] Modern state-of-the-art architectures seek to exploit these "semantic-agnostic" features, which are independent of the image content but highly sensitive to the tampering process.[10]

Artifact Extraction and Forensic Preprocessing Methodologies

A critical prerequisite for effective tampering localization is the suppression of semantic content to reveal forensic artifacts. Standard convolutional neural networks (CNNs) trained for

object recognition tend to focus on high-level shapes and textures, which can be misleading in forensics. Specialized preprocessing layers are therefore employed to highlight low-level signal inconsistencies.

Spatial Rich Model (SRM) and Noise Residual Analysis

The Spatial Rich Model (SRM) represents a collection of handcrafted high-pass filters originally developed for steganalysis. In image forensics, SRM kernels are utilized to extract noise residuals by calculating the difference between a pixel and its neighbors.[12] These residuals capture the local statistical distributions that are often altered during splicing or inpainting. For instance, a common SRM kernel used in the first layer of forensic networks is designed to capture the second-order horizontal or vertical residuals.[12]

The mathematical utility of SRM lies in its ability to act as a regularizer, suppressing the high-amplitude semantic information while amplifying the low-amplitude noise.[12] Research indicates that initializing the initial layers of a CNN with SRM weights significantly boosts localization accuracy compared to random initialization, as it provides the model with a strong prior for noise-aware feature extraction.[13]

Constrained Convolution and BayarConv

While SRM filters are fixed, Bayar Convolution (BayarConv) introduces a learnable mechanism for content suppression. BayarConv imposes a strict constraint on the convolutional kernels during training. Specifically, for a kernel W of size $K \times K$, the weights are constrained such that the central weight $W_{0,0}$ is set to -1, and the sum of all other weights equals 1 [5]:

$$\sum_{i,j \neq 0,0} W_{i,j} = 1, \quad W_{0,0} = -1$$

This constraint ensures that the filter behaves as a high-pass operator, effectively learning the "prediction error" for each pixel based on its neighborhood. By allowing the network to learn the optimal high-pass coefficients, BayarConv can adapt to the specific artifact distributions of the training data, often outperforming fixed SRM filters in detecting sophisticated manipulations like deepfake-based inpainting.[16]

Frequency Domain and Compression Artifacts

Beyond spatial residuals, the frequency domain offers critical forensic clues. Most digital images are stored in the JPEG format, which involves a Discrete Cosine Transform (DCT) and quantization. When an image is tampered with and re-saved, the manipulated region often exhibits a different compression history than the authentic background, leading to "double compression" artifacts.[3] Architectures like CAT-Net and ObjectFormer leverage this by

processing the DCT coefficients directly, allowing the model to detect inconsistencies in the quantization tables that are invisible in the RGB domain.[11]

Artifact Type	Extraction Mechanism	Domain	Forensic Targeted Inconsistency
Noise Residual	SRM Filters	Spatial	Local variance and noise floor mismatches [16]
High-Pass Trace	BayarConv	Spatial	Adaptive prediction errors at boundaries [11]
Compression	DCT Spectrum	Frequency	Mismatched JPEG quantization/double compression [18]
Pixel Dependency	Masked Self-Attention	Spatial	Disruptions in demosaicing/CFA patterns [7]
Local Anomaly	Noiseprint++	Spatial	Deviations from camera-specific ISP fingerprints [20]

Advanced Architectures for Detection and Localization

The landscape of 2024 and 2025 shows a decisive shift toward multi-modal and multi-scale architectures that bridge the gap between pixel-level artifacts and global semantic consistency.

Dual-Branch and Multi-View Networks

A common paradigm in SOTA models is the dual-branch architecture, where one stream processes the RGB image for semantic clues (e.g., edge inconsistencies or lighting mismatches) and a second stream processes forensic artifacts (e.g., SRM noise or DCT coefficients). The MVSS-Net and its enhanced version, MVSS-Net++, are prototypical of this approach. MVSS-Net++ utilizes an Edge-Supervised Branch (ESB) to trace boundary artifacts and a Noise-Sensitive Branch (NSB) to identify noise distribution discrepancies.[10] This

multi-view learning ensures the model is both sensitive to local tampering and specific enough to prevent false alarms on authentic images.[10]

The REFORGE network further evolves this by combining a classification branch for image-level detection with a segmentation branch for pixel-level localization.[22] REFORGE is particularly innovative due to its integration of a reinforcement learning (RL) agent. After the segmentation branch produces an initial mask, the RL agent iteratively refines the mask by treating each pixel as an agent that seeks to maximize a reward based on localization accuracy.[22]

Transformer-Based Forensic Models

The adoption of Transformers, particularly the Vision Transformer (ViT) and SegFormer, has addressed the limitation of CNNs in capturing long-range dependencies. In copy-move forgery, the source and target regions may be far apart; Transformers, through self-attention, can model these distant relationships more effectively than local convolutional kernels.[18]

SegFormer is highly regarded in the forensics community for its hierarchical structure, which outputs multi-scale features, and its lightweight "all-MLP" decoder.[2] For an internship project on Google Colab, SegFormer-B0 or B1 offers an ideal balance between performance and computational efficiency. Unlike standard ViTs that maintain a constant resolution, SegFormer's hierarchical approach allows for precise localization of small tampered regions while maintaining global context.[2] TruFor, another SOTA framework, uses a SegFormer-based encoder to fuse RGB features with Noiseprint++ (a learned camera fingerprint), producing not only a localization map but also a reliability map that quantifies the model's confidence in its predictions.[8]

Explainable and Multi-Modal Models

A rising trend is the development of "explainable" IMDL. FakeShield, for example, is a multimodal framework that combines visual and language models to provide a textual basis for its forensic judgment.[28] It evaluates image authenticity and generates masks while providing a descriptive basis—such as identifying specific textural or lighting inconsistencies—that can be used as evidence in forensic investigations.[28]

Optimization of Loss Functions for Class Imbalance

One of the most significant technical hurdles in image tampering localization is the extreme class imbalance within the dataset. Typically, only a small percentage of pixels (often less than 5%) in a tampered image are actually modified.[29] Standard Binary Cross-Entropy (BCE) loss treats every pixel with equal weight, which leads to models that are biased toward the majority "authentic" class, often resulting in "empty" masks with high overall accuracy but near-zero localization precision.[2]

Dice Loss and the F1 Optimization

Dice Loss, based on the Sorensen-Dice coefficient, is a region-based loss function that directly optimizes for the overlap between the predicted mask (P) and the ground truth (G).[2] It is inherently robust to class imbalance because the background pixels (the majority class) do not contribute to the numerator of the coefficient [31]:

$$L_{Dice} = 1 - \frac{2\sum P \cdot G}{\sum P + \sum G}$$

By focusing only on the intersection of the tampered pixels, Dice Loss forces the model to prioritize the minority class.[32]

Focal Loss and Hard Example Mining

Focal Loss addresses imbalance by reshuffling the importance of examples based on how "hard" they are for the model to classify. It adds a focusing parameter γ to the standard BCE loss, which down-weights the loss contributed by "easy" (confidently predicted) background pixels.[31]

$$L_{Focal} = -(1 - p_t)^{\gamma} \log(p_t)$$

In forensic tasks, pixels at the boundaries of forgeries or in regions with similar textures are "hard" examples. Focal Loss ensures that the gradient updates are dominated by these critical regions rather than the vast authentic background.[31]

Hybrid and Multi-Scale Loss Formulations

The most robust SOTA models utilize a weighted combination of losses. A common configuration includes BCE for global distribution matching, Dice for region overlap, and a specialized Edge Loss to supervise the boundaries.[10]

$$L_{total} = \lambda_1 L_{BCE} + \lambda_2 L_{Dice} + \lambda_3 L_{Edge}$$

Edge loss is often implemented by calculating the binary cross-entropy between the predicted edges and a ground truth edge map (derived by dilating and eroding the tampering mask).[10] This ensures that the model learns to produce sharp, distinct boundaries rather than blurry, ambiguous localization maps.

Performance Optimization for Google Colab T4 GPUs

The NVIDIA T4 GPU provided by Google Colab is an efficient accelerator with 16 GB of VRAM, optimized primarily for inference but capable of training mid-sized forensic models.[37] To maximize its utility for an internship assignment, several memory and compute optimizations must be applied.

Mixed Precision Training (AMP)

Mixed precision training involves performing the majority of forward and backward pass computations in 16-bit floating point (FP16) while maintaining weight updates in 32-bit (FP32).[39] This technique leverages the T4's Tensor Cores, leading to a 1.5x to 2x speedup in training and a reduction in memory usage by approximately 50%.[39] PyTorch's autocast and GradScaler modules are the industry-standard tools for implementing this, ensuring numerical stability by scaling gradients to prevent underflow.[39]

Gradient Checkpointing and Accumulation

For deeper architectures like ObjectFormer or ProFact, the activation maps can exceed the 16 GB VRAM limit. Gradient checkpointing is a technique that reduces memory consumption by trading compute for memory; instead of storing all intermediate activations for the backward pass, it recomputes them on-the-fly.[43] This allows for the use of larger batch sizes or higher input resolutions (e.g., 512x512) which are critical for capturing high-frequency forensic artifacts.[44]

If the desired batch size still does not fit in memory, gradient accumulation can be used. This involves accumulating gradients over multiple small "micro-batches" before performing a single optimizer update, effectively simulating a larger batch size without increasing the peak VRAM load.[46]

Memory-Efficient Optimizers and Quantization

Modern optimizers like AdamW can be memory-intensive due to the storage of two moments per parameter. Using 8-bit optimizers (e.g., via the bitsandbytes library) can reduce the optimizer's memory footprint by up to 75%.[46] Furthermore, if using a large pre-trained backbone, Low-Rank Adaptation (LoRA) can be employed to fine-tune only a small fraction of the parameters, significantly reducing the memory required for gradients and optimizer states.[46]

Optimization Technique	Mechanism	Impact on T4 Performance
Mixed Precision (AMP)	FP16/FP32 hybrid math	~2x faster training; 50%

		VRAM reduction [39]
Gradient Checkpointing	Recomputes activations	Enables 2x-3x larger batch sizes [44]
8-bit AdamW	Quantizes optimizer states	75% reduction in optimizer memory [46]
LoRA / QLoRA	Parameter-efficient tuning	Fine-tunes large models on <8GB VRAM [46]
FlashAttention	Reordered attention math	Faster transformer training; less memory [49]

Benchmarking and Evaluation Methodologies

Evaluating an image forgery detector requires more than simple accuracy metrics. The model must be assessed for its localization precision and its robustness against common "anti-forensic" attacks designed to hide tampering traces.

Evaluation Metrics

1. **Pixel-Level F1-Score and IoU:** These are the primary metrics for localization, measuring the harmony between precision and recall at the pixel level.[51]
2. **Image-Level Detection Accuracy:** Measures the model's ability to correctly classify an image as "authentic" or "tampered".[4]
3. **Area Under the Curve (AUC-ROC):** Provides a threshold-independent measure of the model's ability to distinguish between forged and authentic pixels.[29]
4. **Matthews Correlation Coefficient (MCC):** Useful for highly imbalanced datasets, as it accounts for all four quadrants of the confusion matrix (TP, TN, FP, FN).

The Calibration Gap and Oracle-F1

Recent benchmarks have identified a "Calibration Gap" in forensic models. While many models achieve high AUC-ROC scores, their F1-scores at a fixed threshold (e.g., $\tau = 0.5$) are often near-zero.[29] This is because the probability scores are not well-calibrated for the extreme rarity of tampered pixels. To address this, "Oracle-F1"—the best achievable F1-score across all possible thresholds—is used as a secondary benchmark to evaluate the model's discriminative potential independently of threshold selection.[29]

Dataset	Metric	SOTA Model (e.g.,	Performance

		MVSS / TruFor)	Value
CASIA v2.0	Pixel-F1	VAAS (2025)	94.1% [52]
CASIA v2.0	Pixel-IoU	VASLNet (2024)	85.1% [52]
COVERAGE	Pixel-F1	MMFD-Net (2024)	50.4% [51]
COVERAGE	Pixel-F1	MVSS-Net++ (2022)	48.2% [51]
CASIA v1.0	Pixel-F1	IF-OSN (2022)	68.6% [54]
NIST16	Pixel-F1	DAE-Net (2024)	87.1% [53]
IMD2020	Pixel-F1	DAE-Net (2024)	33.8% [53]

Robustness Evaluation

A "bonus" requirement for high-tier forensic models is robustness against distortions. Evaluation should be performed on a test set that has been subjected to:

- **JPEG Compression:** Testing at QF levels 50, 70, and 90.[10]
- **Resizing and Scaling:** Bilinear and bicubic resizing between 0.5x and 2.0x.[10]
- **Noise Addition:** Gaussian white noise with varying variance.[55]
- **Screenshot/Re-capturing:** Simulating the artifacts introduced when an image is displayed and then re-photographed or screen-captured.[10]

Strategic Research and Implementation Roadmap

For the internship assignment, a phased approach is recommended to ensure both a working baseline and a competitive, SOTA-aligned final model.

Phase 1: Data Infrastructure and Preprocessing

The focus should be on the CASIA v2.0 and COVERAGE datasets, which provide a robust mix of splicing and copy-move forgeries.[4]

- **Pipeline Development:** Use torchvision and albumentations to build a data loader that performs online augmentation, including rotation, flipping, and random JPEG compression.[10]
- **Artifact Channel:** Implement an SRM filter layer as the first operation. A multi-channel output representing different noise residuals should be concatenated with the RGB input.[12]

- **Mask Preparation:** Ensure that ground truth masks are binary and aligned. For copy-move, specifically distinguish between source and target regions if possible (e.g., using BusterNet principles).[6]

Phase 2: Architecture Selection and Baseline Training

Given the Colab T4 constraints, a hierarchical transformer or a dual-stream CNN is the optimal choice.

- **Backbone:** Select a SegFormer-B1 encoder pre-trained on ADE20K or ImageNet. This provides a strong starting point for semantic understanding.[2]
- **Fusion Strategy:** Implement a dual-stream encoder where the first stream is the standard SegFormer (RGB) and the second stream is a lightweight CNN (Noise Residuals). Features should be fused using a Spatial Attention module to highlight local inconsistencies.[10]
- **Initial Training:** Train with a combination of BCE and Dice Loss. Use a small batch size (e.g., 4 or 8) with gradient accumulation to fit within the 16 GB VRAM.[2]

Phase 3: Forensic Refinement and Multi-Task Learning

Enhance the model by adding specialized forensic supervisory signals.

- **Edge Supervision:** Implement an auxiliary branch that predicts the boundaries of tampered regions. Use a multi-scale edge loss to ensure the model focuses on the delicate artifacts found at the periphery of spliced objects.[10]
- **Optimization:** Enable Mixed Precision (AMP) to accelerate training. Profile the memory usage and, if possible, increase the input resolution to 512x512 to preserve high-frequency artifacts.[39]
- **Refinement:** Implement a coarse-to-fine strategy where the model first predicts a low-resolution mask and then refines it using a dedicated refinement block (e.g., as in PSCC-Net or ProFact).[18]

Phase 4: Evaluation and Documentation

The final deliverable must align with the rigorous testing requirements of the assignment.[4]

- **Benchmarking:** Evaluate the final model on the CASIA v2.0 test set and report Pixel-F1, Pixel-IoU, and Image-Accuracy.[4]
- **Visual Documentation:** Generate a set of comprehensive visualizations for the report:
  1. Original Image.
  2. Ground Truth Mask.
  3. Predicted Localization Map (Probability Heatmap).
  4. Overlay Visualization (Mask on Image).
  5. Confidence/Reliability Map (if using a TruFor-style decoder).[4]
- **Robustness Report:** Create a table showing how performance degrades under JPEG compression (QF=50) and Resizing (0.5x), fulfilling the bonus point criteria.[4]

Mechanism of Pixel Correlation Disruption in Forensics

A deeper insight into why these models work lies in the demosaicing process. Digital cameras use a Color Filter Array (CFA), most commonly a Bayer filter, to capture color information. Because each pixel only records one of three colors (Red, Green, or Blue), the camera's Image Signal Processor (ISP) must interpolate the missing colors using the surrounding pixels. This creates a specific, periodic correlation between neighboring pixels.[7]

When an image is tampered with, these periodic correlations are disrupted. A spliced object from a different camera will have a different CFA pattern or demosaicing algorithm, creating a statistical mismatch at the boundary.[7] Advanced forensic models, such as those employing masked self-attention or constrained convolutions, are designed specifically to detect these low-level geometric inconsistencies. This is why the SRM and BayarConv preprocessing steps are so effective; they strip away the high-level semantic content (which can be easily faked by generative AI) and focus on the fundamental mathematical structure of the digital signal.[7]

Future Outlook and Generative AI Challenges

The rapid advancement of generative AI (AIGC) presents a dual-edged sword for image forensics. On one hand, diffusion-based inpainting can create forgeries that are visually perfect and preserve the global semantic logic of a scene.[61] On the other hand, these generative processes often leave distinct "generative noise" patterns that differ significantly from camera-sensor noise.[61]

SOTA research in 2025 is increasingly focusing on "Zero-Shot" localization and adaptation to novel generative models. Techniques such as using the internal generative process of stable diffusion models to reveal forgery artifacts (as seen in the CLUE framework) represent the next step in this arms race.[61] For the current assignment, the focus remains on capturing the transition between the traditional artifacts (splicing/copy-move) and the emerging generative traces, ensuring a robust and versatile forensic toolset.

By adhering to this strategic framework, an intern can develop a model that not only meets the assignment criteria but aligns with the cutting-edge methodologies of the professional forensics community. The combination of lightweight SegFormer backbones, artifact-suppressing preprocessing, and robust optimization for class-imbalanced data provides a technically sound and high-performing solution for the task of image tampering detection and localization.

Works cited

1. A comprehensive review of deep learning techniques for image tampering detection, accessed March 10, 2026, https://www.researchgate.net/publication/398245250_A_comprehensive_review_

of_deep_learning_techniques_for_image_tampering_detection
2. Hierarchically Structured Transformer Encoders: Image Forgery ..., accessed March 10, 2026, https://medium.com/@jsilvawasd/hierarchically-structured-transformer-encoders-document-forgery-semantic-segmentation-5793c2bf3ec1
3. Tampering Detection and Segmentation Model for Multimedia Forensic - The Science and Information (SAI) Organization, accessed March 10, 2026, https://thesai.org/Downloads/Volume14No9/Paper_92-Tampering_Detection_and_Segmentation_Model.pdf
4. Internship Assignment_ Tampered Image Detection & Localization.pdf
5. Can We Get Rid of Handcrafted Feature Extractors? SparseViT: Nonsemantics-Centered, Parameter-Efficient Image Manipulation Localization Through Spare-Coding Transformer - arXiv, accessed March 10, 2026, https://arxiv.org/html/2412.14598v1
6. Enhancing Copy-Move Forgery Detection via Attention-Based Similarity Modeling - World Scientific Publishing, accessed March 10, 2026, https://www.worldscientific.com/doi/pdf/10.1142/S0218001425500442?download=true
7. Pixel-Inconsistency Modeling for Image Manipulation Localization - arXiv.org, accessed March 10, 2026, https://arxiv.org/html/2310.00234v2
8. TruFor: Leveraging All-Round Clues for ... - CVF Open Access, accessed March 10, 2026, https://openaccess.thecvf.com/content/CVPR2023/papers/Guillaro_TruFor_Leveraging_All-Round_Clues_for_Trustworthy_Image_Forgery_Detection_and_CVPR_2023_paper.pdf
9. Forged Copy-Move Recognition Using Convolutional Neural Network - Al-Nahrain Journal of Science (ANJS), accessed March 10, 2026, https://anjs.edu.iq/index.php/anjs/article/download/2335/1837/4518
10. MVSS-Net: Multi-View Multi-Scale Supervised Networks for Image ..., accessed March 10, 2026, https://pubmed.ncbi.nlm.nih.gov/35671312/
11. Can We Get Rid of Handcrafted Feature Extractors? SparseViT: Nonsemantics-Centered, Parameter-Efficient Image Manipulation Localization through Spare-Coding Transformer - arXiv.org, accessed March 10, 2026, https://arxiv.org/html/2412.14598v2
12. Multitask Image Splicing Tampering Detection Based on Attention Mechanism, accessed March 10, 2026, https://www.researchgate.net/publication/364464845_Multitask_Image_Splicing_Tampering_Detection_Based_on_Attention_Mechanism
13. Data-Dependent Scaling of CNN's First Layer for Improved Image Manipulation Detection - Gipsa-lab, accessed March 10, 2026, https://www.gipsa-lab.grenoble-inp.fr/~kai.wang/papers/IWDW20.pdf
14. Learning Rich Features for Image Manipulation Detection - CVF Open Access, accessed March 10, 2026, https://openaccess.thecvf.com/content_cvpr_2018/papers/Zhou_Learning_Rich_Features_CVPR_2018_paper.pdf

15. Image Forgery Detection. Introduction | by Kishan | Medium, accessed March 10, 2026, https://kishanraj-16649.medium.com/image-forgery-detection-ae7a8101ddf
16. MMFusion: Combining Image Forensic Filters for Visual Manipulation Detection and Localization - arXiv.org, accessed March 10, 2026, https://arxiv.org/html/2312.01790v2
17. Exploring Multi-Modal Fusion for Image Manipulation Detection and Localization - arXiv, accessed March 10, 2026, https://arxiv.org/html/2312.01790v1
18. DocForge-Bench: A Comprehensive Benchmark for Document Forgery Detection and Analysis - arXiv, accessed March 10, 2026, https://arxiv.org/html/2603.01433v1
19. ObjectFormer for Image Manipulation Detection and Localization ..., accessed March 10, 2026, https://www.researchgate.net/publication/363909282_ObjectFormer_for_Image_Manipulation_Detection_and_Localization
20. TruFor: Leveraging All-Round Clues for Trustworthy Image Forgery Detection and Localization - grip-unina, accessed March 10, 2026, https://grip-unina.github.io/TruFor/
21. MVSS-Net: Multi-View Multi-Scale Supervised Networks for Image Manipulation Detection, accessed March 10, 2026, https://www.researchgate.net/publication/357114617_MVSS-Net_Multi-View_Multi-Scale_Supervised_Networks_for_Image_Manipulation_Detection
22. REFORGE: A Robust Ensemble for Image Forgery Detection and Localization in Social Network Images - IEEE Xplore, accessed March 10, 2026, https://ieeexplore.ieee.org/iel8/6287639/11323511/11346498.pdf
23. (PDF) REFORGE: A Robust Ensemble for Image Forgery Detection and Localization in Social Network Images - ResearchGate, accessed March 10, 2026, https://www.researchgate.net/publication/399712748_REFORGE_A_Robust_Ensemble_for_Image_Forgery_Detection_and_Localization_in_Social_Network_Images
24. Enhanced Copy-Move Forgery Detection in Images Using Hybrid Deep Learning Approach - EPJ Web of Conferences, accessed March 10, 2026, https://www.epj-conferences.org/articles/epjconf/pdf/2025/26/epjconf_icatcict2025_01042.pdf
25. SegFormer - NVIDIA Docs, accessed March 10, 2026, https://docs.nvidia.com/tao/tao-toolkit-archive/5.2.0/text/semantic_segmentation/segformer.html
26. Research on the performance of the SegFormer model with fusion of edge feature extraction for metal corrosion detection - PMC, accessed March 10, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC11890780/
27. Image Forgery Detection - grip-unina, accessed March 10, 2026, https://www.grip.unina.it/multimedia-forensics/image-forgery-detection
28. FakeShield: Explainable Image Forgery Detection and Localization via Multi-modal Large Language Models | OpenReview, accessed March 10, 2026, https://openreview.net/forum?id=pAQzEY7M03
29. DOCFORGE-BENCH: A Comprehensive Benchmark for Document Forgery

Detection and Analysis - arXiv, accessed March 10, 2026, https://arxiv.org/pdf/2603.01433
30. (PDF) DOCFORGE-BENCH: A Comprehensive Benchmark for Document Forgery Detection and Analysis - ResearchGate, accessed March 10, 2026, https://www.researchgate.net/publication/401469940_DOCFORGE-BENCH_A_Comprehensive_Benchmark_for_Document_Forgery_Detection_and_Analysis
31. Instance segmentation loss functions - SoftwareMill, accessed March 10, 2026, https://softwaremill.com/instance-segmentation-loss-functions/
32. Understanding Loss Functions for Deep Learning Segmentation Models - Medium, accessed March 10, 2026, https://medium.com/@devanshipratiher/understanding-loss-functions-for-deep-learning-segmentation-models-30187836b30a
33. The combined focal loss and dice loss function improves the segmentation of beta-sheets in medium-resolution cryo-electron-microscopy density maps - PMC, accessed March 10, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC11590252/
34. The Loss Functions That Actually Matter in 2025 | by Pranav Prakash I GenAI I AI/ML I DevOps I | Medium, accessed March 10, 2026, https://medium.com/@pranavprakash4777/the-loss-functions-that-actually-matter-in-2025-41b044b2645e
35. Progressive Feedback-Enhanced Transformer for Image Forgery Localization - arXiv, accessed March 10, 2026, https://arxiv.org/html/2311.08910v2
36. Multi-scale and deeply supervised network for image splicing localization - Frontiers, accessed March 10, 2026, https://www.frontiersin.org/journals/artificial-intelligence/articles/10.3389/frai.2025.1655073/full
37. arch.ipynb - Colab, accessed March 10, 2026, https://colab.research.google.com/github/d2l-ai/d2l-tvm-colab/blob/master/chapter_gpu_schedules/arch.ipynb
38. Understanding Google Colab Free GPU in detail | by Mehul Gupta | Data Science in Your Pocket | Medium, accessed March 10, 2026, https://medium.com/data-science-in-your-pocket/understanding-google-colab-free-gpu-in-detail-15074081d494
39. How to implement mixed precision training with PyTorch? - Tencent Cloud, accessed March 10, 2026, https://www.tencentcloud.com/techpedia/126165
40. NVIDIA Apex: Tools for Easy Mixed-Precision Training in PyTorch | NVIDIA Technical Blog, accessed March 10, 2026, https://developer.nvidia.com/blog/apex-pytorch-easy-mixed-precision-training/
41. What Every User Should Know About Mixed Precision Training in pytorch | by Hey Amit | Data Scientist's Diary | Medium, accessed March 10, 2026, https://medium.com/data-scientists-diary/what-every-user-should-know-about-mixed-precision-training-in-pytorch-63c6544e5a05
42. Mixed Precision Training | Explanation and PyTorch Implementation from Scratch - YouTube, accessed March 10, 2026, https://www.youtube.com/watch?v=hHpC9Sywh4U

43. Efficient Llama Training with Gradient Checkpointing and Adapters - Colab, accessed March 10, 2026, https://colab.research.google.com/github/Adapter-Hub/adapters/blob/main/notebooks/Gradient_Checkpointing_Llama.ipynb
44. How to implement gradient checkpointing for transformer models to reduce memory usage?, accessed March 10, 2026, https://massedcompute.com/faq-answers/?question=How%20to%20implement%20gradient%20checkpointing%20for%20transformer%20models%20to%20reduce%20memory%20usage?
45. Memory consumption qlora with gradient checkpointing - Hugging Face Forums, accessed March 10, 2026, https://discuss.huggingface.co/t/memory-consumption-qlora-with-gradient-checkpointing/70948
46. Fine-Tuning Code GenAI model on Google Colab T4 GPU: A Step-by-Step Guide, accessed March 10, 2026, https://www.tothenew.com/blog/fine-tuning-code-genai-model-on-google-colab-t4-gpu-a-step-by-step-guide/
47. Understanding the Limits of Google Colab Pro RAM and GPU Usage - YouTube, accessed March 10, 2026, https://www.youtube.com/watch?v=GY7yObl6Zi4
48. google-research/vision_transformer - GitHub, accessed March 10, 2026, https://github.com/google-research/vision_transformer
49. A practical guide to GPU memory for fine-tuning AI models | Google Cloud Blog, accessed March 10, 2026, https://cloud.google.com/blog/topics/developers-practitioners/decoding-high-bandwidth-memory-a-practical-guide-to-gpu-memory-for-fine-tuning-ai-models/
50. Introducing Mixed Precision Training in Opacus - PyTorch, accessed March 10, 2026, https://pytorch.org/blog/introducing-mixed-precision-training-in-opacus/
51. MMFD-Net: A Novel Network for Image Forgery Detection and ..., accessed March 10, 2026, https://www.mdpi.com/2227-7390/13/19/3136
52. VAAS: Vision-Attention Anomaly Scoring for Image Manipulation Detection in Digital Forensics - arXiv.org, accessed March 10, 2026, https://arxiv.org/html/2512.15512v1
53. DAE-Net: Dual Attention Mechanism and Edge Supervision Network for Image Manipulation Detection and Localization - IEEE Xplore, accessed March 10, 2026, https://ieeexplore.ieee.org/iel8/19/4407674/10658981.pdf
54. Weakly-Supervised Image Forgery Localization via Vision-Language Collaborative Reasoning Framework - arXiv, accessed March 10, 2026, https://arxiv.org/html/2508.01338v1
55. Robust Detection and Localization of Image Copy-Move Forgery Using Multi-Feature Fusion, accessed March 10, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC12941880/
56. image_tampering_detection/phase2 (SRM+Fake,ELA+Unet).ipynb at master - GitHub, accessed March 10, 2026, https://github.com/enviz/image_tampering_detection/blob/master/phase2%20(SRM%2BFake%2CELA%2BUnet).ipynb

57. SegFormer Tutorial: Master Semantic Segmentation Fast - Labellerr, accessed March 10, 2026, https://www.labellerr.com/blog/segformer/
58. (PDF) Effective Image Tampering Localization via Enhanced Transformer and Co-Attention Fusion - ResearchGate, accessed March 10, 2026, https://www.researchgate.net/publication/379266243_Effective_Image_Tampering_Localization_via_Enhanced_Transformer_and_Co-Attention_Fusion
59. TruFor: Leveraging all-round clues for trustworthy image forgery detection and localization, accessed March 10, 2026, https://www.researchgate.net/publication/366497579_TruFor_Leveraging_all-round_clues_for_trustworthy_image_forgery_detection_and_localization
60. Convolutional PyTorch debayering / demosaicing layers - GitHub, accessed March 10, 2026, https://github.com/cheind/pytorch-debayer
61. CLUE: Leveraging Low-Rank Adaptation to Capture Latent Uncovered Evidence for Image Forgery Localization - arXiv, accessed March 10, 2026, https://arxiv.org/html/2508.07413v1
62. SDiFL: Stable Diffusion-Driven Framework for Image Forgery Localization - arXiv, accessed March 10, 2026, https://arxiv.org/pdf/2508.20182
63. CVPR Poster Towards Enhanced Image Inpainting: Mitigating Unwanted Object Insertion and Preserving Color Consistency, accessed March 10, 2026, https://cvpr.thecvf.com/virtual/2025/poster/35187
64. No Pixel Left Behind: A Detail-Preserving Architecture for Robust High-Resolution AI-Generated Image Detection | OpenReview, accessed March 10, 2026, https://openreview.net/forum?id=9QQ3Kc2hj6